



• PRIVACYPAL FOR ENTERPRISE

Govern every AI. Without blocking a single user.

The governance and enablement layer for AI — so your enterprise can adopt AI everywhere it creates value, with trust and compliance built in. Not a brake on innovation. The layer that unlocks it.

[Book a demo →](#)

[Watch the product intro](#)

```
# Install on any desktop or server
$ curl -sSL https://get.privacypal.ai
| sh

# Every AI agent on this machine is
# now protected by PrivacyPal Cloud.

$ privacypal status
✓ Gateway running on 127.0.0.1:7878
✓ Watching 4 local AI processes
✓ Routing via
  privacypal.yourco.internal
✓ 0 prompts left this network
```

Your people are already all-in on AI. Your governance **isn't.**

01 **AI is already ~60% of the workday.**

Teams run client data, code and strategy through ChatGPT, Claude and Copilot every day — with or without you.

02 **Block it, and innovation stalls.**

The only lever most security teams have is "no." AI initiatives die in the review queue while competitors ship.

03 **Allow it ungoverned, and trust is on the line.**

Unprotected prompts leak PII, PHI and IP to third-party models — risking trust your firm spent a lifetime earning.

Stop blocking AI. **Start governing it.**

The trust layer beneath your AI operating system.

DATA IN MOTION • GATEWAY

Govern every AI request

Installs at the network edge and intercepts AI traffic with zero code changes. Sensitive values are swapped for Privacy Twins on the way out and restored on return — provider-agnostic, streaming-safe.

OpenAI • Anthropic • Google • Mistral • Llama • Own fine-tune

DATA AT REST • NETWORK DSPM

Govern every data store

The same install continuously discovers, classifies, prioritizes and remediates sensitive data across the AI data layer — databases, lakes, RAG pipelines, vector stores and agent memory.


SQL • PostgreSQL • MongoDB • Supabase • Object • Vector/RAG

One layer. Every AI, every endpoint, every data store — governed and enabled.

• THE MECHANISM • PRIVACY TWINS

Language-based privacy for a language-based threat.

Others stop at redaction — which breaks the model's output and kills adoption. PrivacyPal substitutes a **context-preserving Privacy Twin**, then restores the real values on the way back. The work gets done. The data never leaves.



• THE DEPLOYMENT CHECKLIST

Sovereign by design. Drop in anywhere your workloads live.



Zero code changes

Intercepts AI traffic at the network edge. Existing agents, copilots and SDKs keep their config — every request twin-swapped.



AWS · Azure · GCP · bare metal

Terraform, Helm, or single-container deploy. Production-ready in an afternoon, inside your tenancy — no third-party DPAs.



Air-gapped option

Run detection with on-prem models. No outbound traffic. Built for regulated enterprises — Healthcare, FSI, Defense.



Provider-agnostic gateway

OpenAI / Anthropic / Google compatible. Route to GPT, Claude, Gemini, Mistral, Llama or your own fine-tune.



Streaming-safe

Detects and swaps mid-stream without breaking token-by-token responses. Users never see a hitch.



Horizontal scale & 24/7 SLA

Stateless workers scale to 100k req/s. P99 under half a second. Dedicated SE. 24/7 support included.

Govern the AI data layer like the database layer.



Agentless discovery

Connect a store and it's in the scan queue. SQL, Mongo, Supabase, Drive, object storage — read in place, no agents.



Context-aware classification

The intelligence behind Privacy Twins reads data like a human. PII, PHI, financials, schemas — no regex to maintain.



Risk that's actually risk

Sensitivity tied to access, exposure and business purpose. An SSN in a public bucket is critical; a dev fixture is not.



One-click remediation

Revoke access, mask columns, trigger a workflow, or route findings to the owner with full context. Cloud closes the loop.



RAG, vector & agent memory

Scan the AI data layer the way you scan the database layer — where models read from, and where data quietly leaks in.



Sovereign by design

Scans run inside your network against the same gateway. Findings, classifications and remediation never leave your tenancy.

Every twin and action, anchored on-chain.



Tokenized, never exposed

Sensitive values become Privacy Twins. Only cryptographic hashes and content IDs go on-chain — never the PII.



AI activity, audited on-chain

Each session is one immutable peaq transaction — user, machine identity, model and twin hashes. Tamper-proof.



DID-secured vault access

A peaq DID is each machine's cryptographic identity. An encrypted vault and DID-gated endpoints guard every twin.

PEAQ · PRIVACY_TWIN_BATCH_AUDIT ✓ confirmed

tx 0x9f3c...a847 block #4,182,663

3a7f...e1 → c19e...b4 name

88b2...9c → 4471...42 mrn

d04a...7f → a8e2...10 address

machine DID did:peaq:0x0841...84ef

IPFS CID QmP8Ud...1YGT ✓ pinned

payload hashes only · 0 bytes PII

Stateless. Sovereign. Throughput-grade.



100k
requests / sec

<0.5s
P99 latency

Zero
knowledge · your keys

Air-gap
ready · no egress

Total governance. Full visibility. Provable trust.

Full visibility

- ✓ Every prompt to every model — logged: who, tool, data type, when.
- ✓ Findings, classifications & remediation never leave your tenancy.
- ✓ Immutable on-chain audit trail for high-assurance evidence.

Compliance-ready exports

GDPR

CCPA

HIPAA

SOC 2

ISO 27001

User	Model	Data Type	Status
a.chen	GPT-4o	Customer PII	● Protected
m.ruiz	Claude	Financial	● Protected
j.park	Gemini	Source code	● Protected
s.okoro	Llama (on-prem)	PHI	● Masked + flagged
d.lang	GPT-4o	API keys	● Protected

“

We shipped an AI assistant on patient records in six weeks. Legal signed it in two days. That's PrivacyPal Cloud.

— CTO, healthcare SaaS · Series B

10M+

ELEMENTS GOVERNED IN REAL TIME

Days

NOT QUARTERS, TO LAUNCH



PLUGANDPLAY

• THE GOVERNANCE LAYER FOR AI

Stop blocking AI. Start governing it.

The layer that turns AI from your biggest risk into your biggest advantage.

[Book a demo →](#)

[Watch the product intro](#)

[privacypal.ai](#)

[Book a demo](#)

[Solution Engineering](#)